



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Bagging

Elie Kawerk
Data Scientist



Ensemble Methods

Voting Classifier

- same training set,
- \neq algorithms.

Bagging

- one algorithm,
- \neq subsets of the training set.

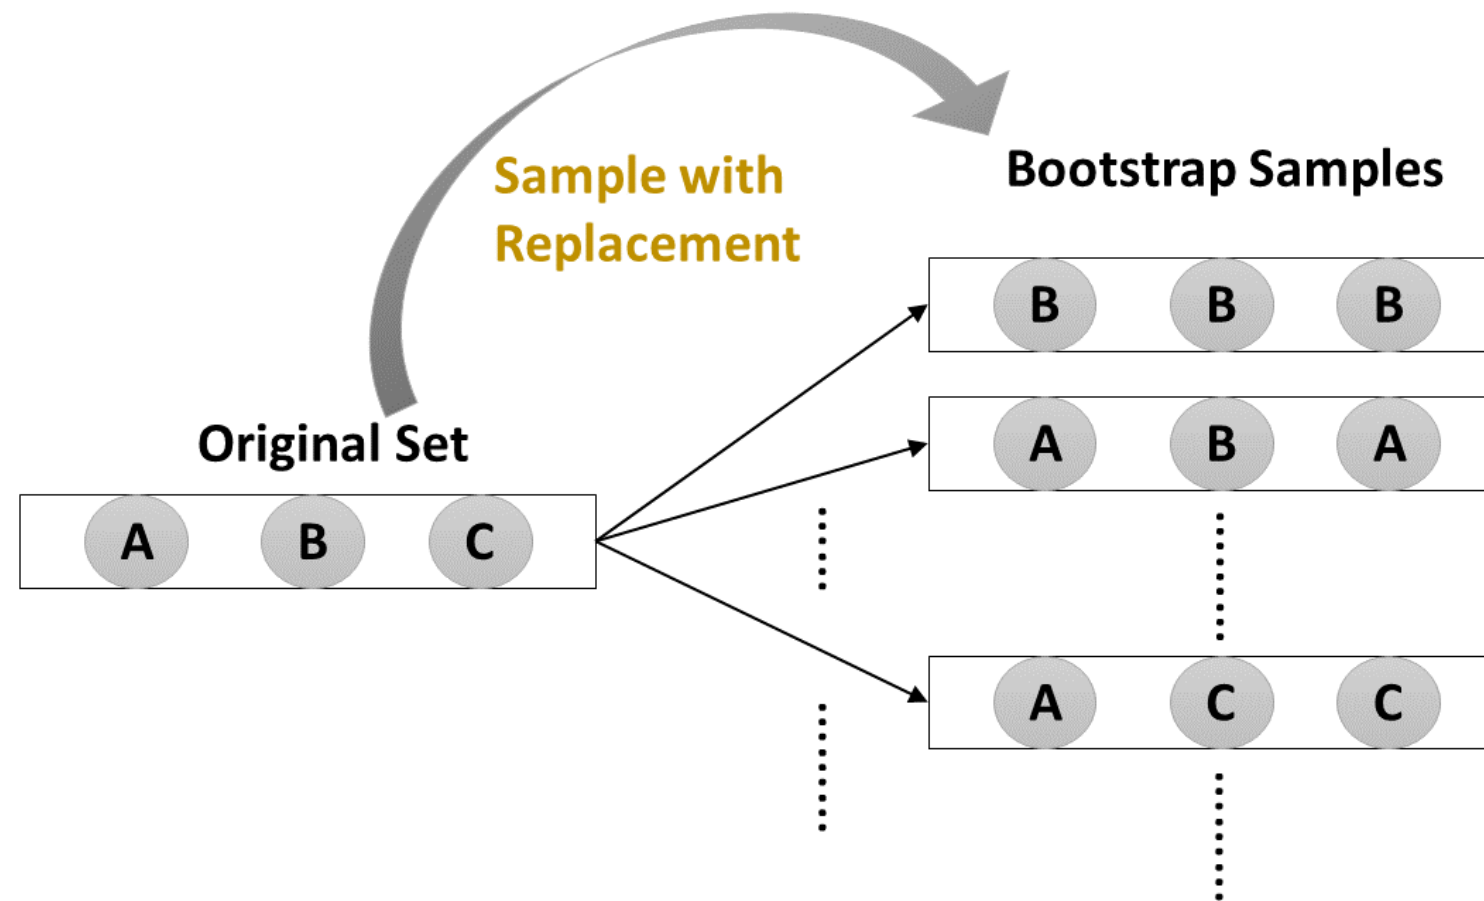


Bagging

- Bagging: Bootstrap Aggregation.
- Uses a technique known as the bootstrap.
- Reduces variance of individual models in the ensemble.

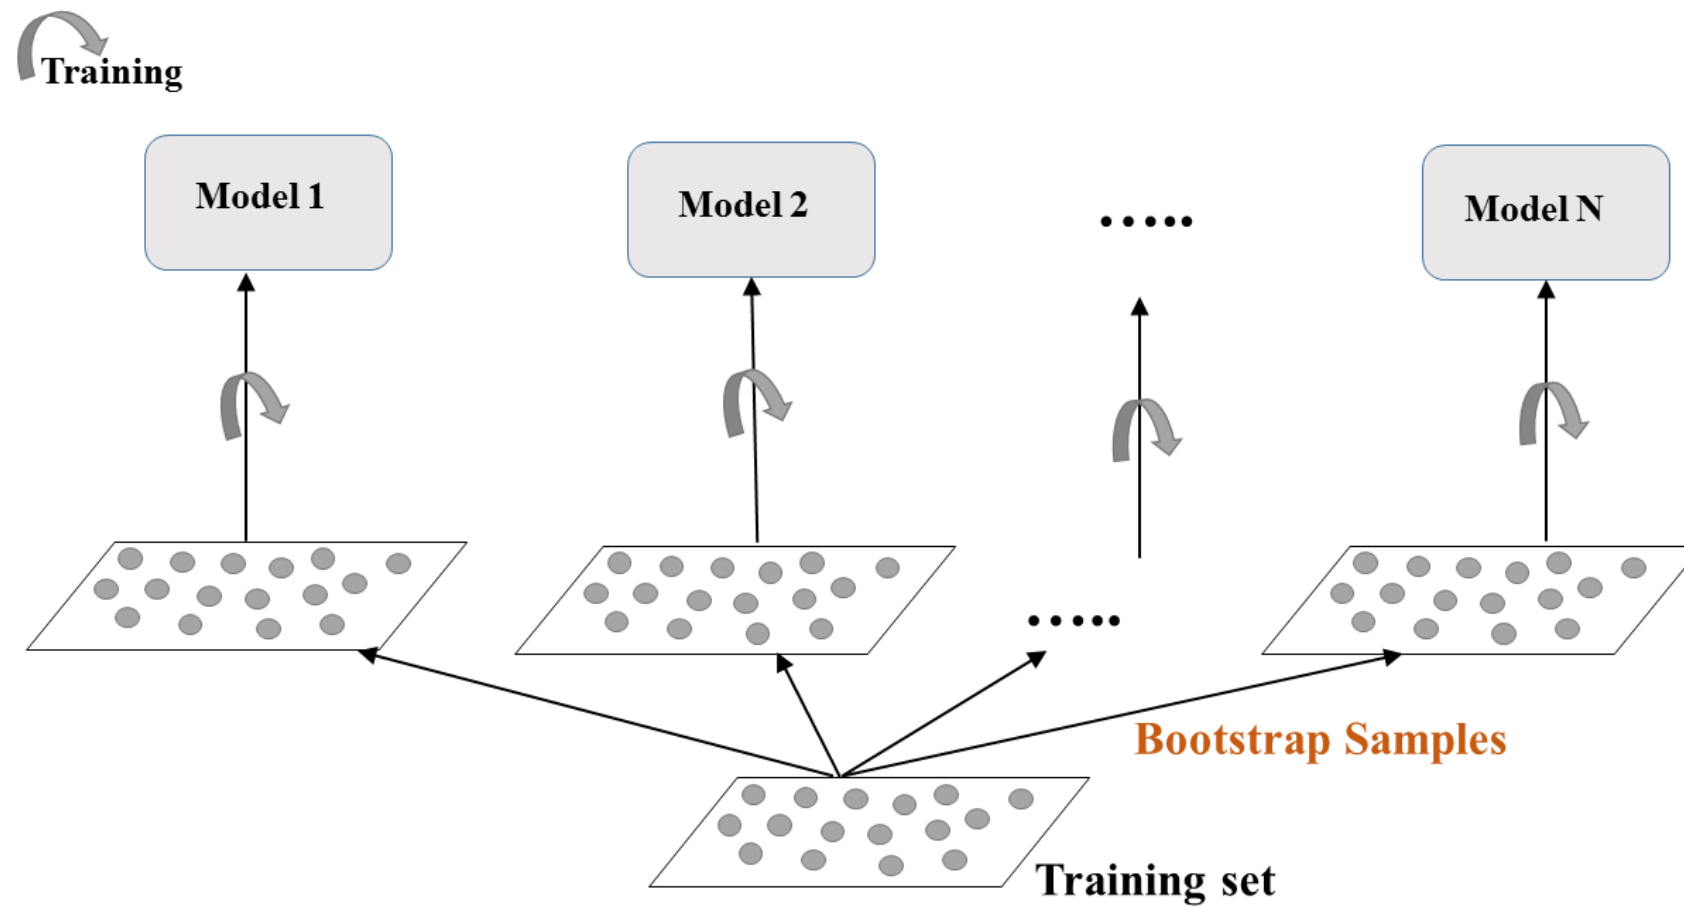


Bootstrap



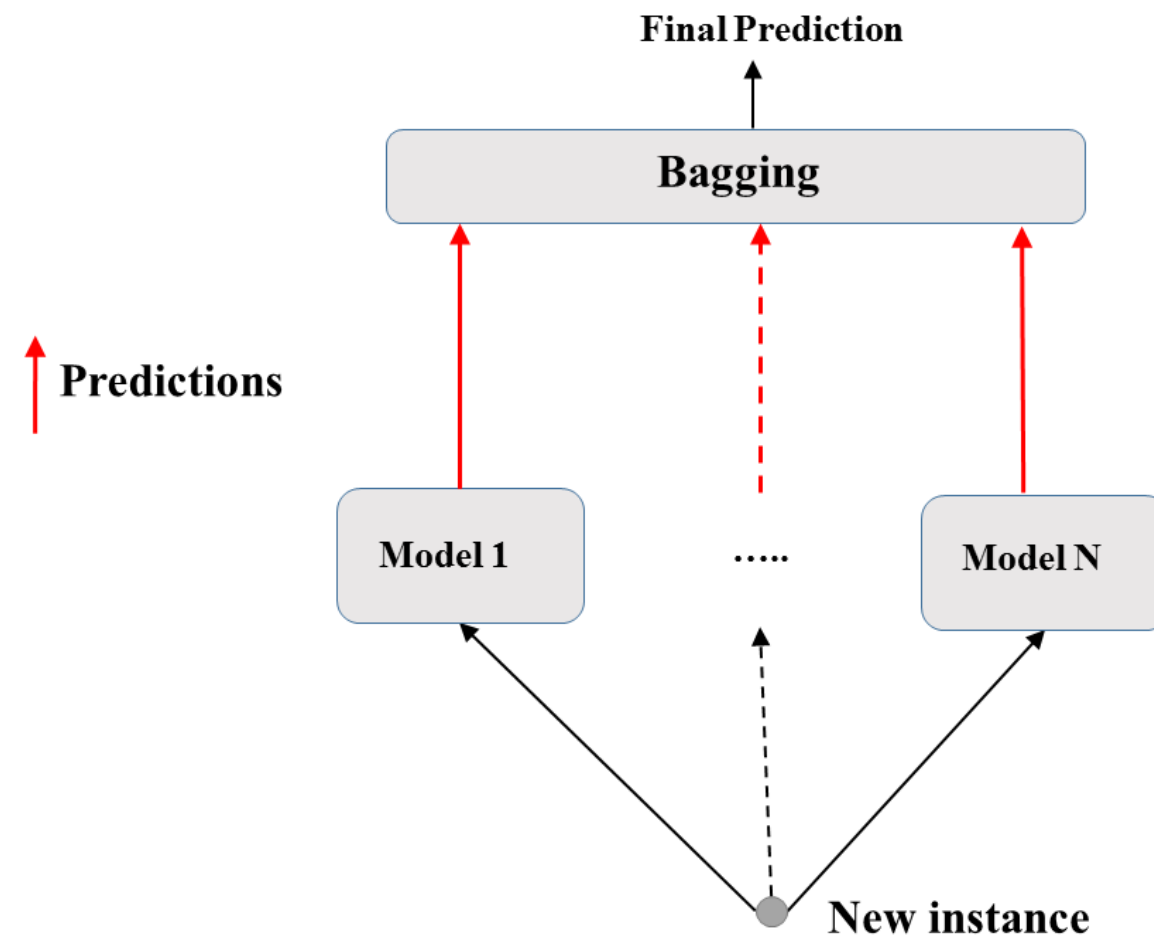


Bagging: Training





Bagging: Prediction





Bagging: Classification & Regression

Classification:

- Aggregates predictions by majority voting.
- `BaggingClassifier` in `scikit-learn`.

Regression:

- Aggregates predictions through averaging.
- `BaggingRegressor` in `scikit-learn`.



Bagging Classifier in sklearn (Breast-Cancer dataset)

```
# Import models and utility functions
In [1]: from sklearn.ensemble import BaggingClassifier
In [2]: from sklearn.tree import DecisionTreeClassifier
In [3]: from sklearn.metrics import accuracy_score
In [4]: from sklearn.model_selection import train_test_split

# Set seed for reproducibility
In [5]: SEED = 1

# Split data into 70% train and 30% test
In [6]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size=0.3,
                        stratify=y,
                        random_state=SEED)
```


Bagging Classifier in sklearn (Breast-Cancer dataset)

```
# Instantiate a classification-tree 'dt'
In [7]: dt = DecisionTreeClassifier(max_depth=4, min_samples_leaf=0.16,
                                     random_state=SEED)

# Instantiate a BaggingClassifier 'bc'
In [8]: bc = BaggingClassifier(base_estimator=dt, n_estimators=300,
                               n_jobs=-1)

# Fit 'bc' to the training set
In [9]: bc.fit(X_train, y_train)

# Predict test set labels
In [10]: y_pred = bc.predict(X_test)

# Evaluate and print test-set accuracy
In [11]: accuracy = accuracy_score(y_test, y_pred)

In [12]: print('Accuracy of Bagging Classifier: {:.3f}'.format(accuracy))

Out[12]: Accuracy of Bagging Classifier: 0.936
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Out Of Bag Evaluation

Elie Kawerk
Data Scientist



Bagging

- some instances may be sampled several times for one model,
- other instances may not be sampled at all.

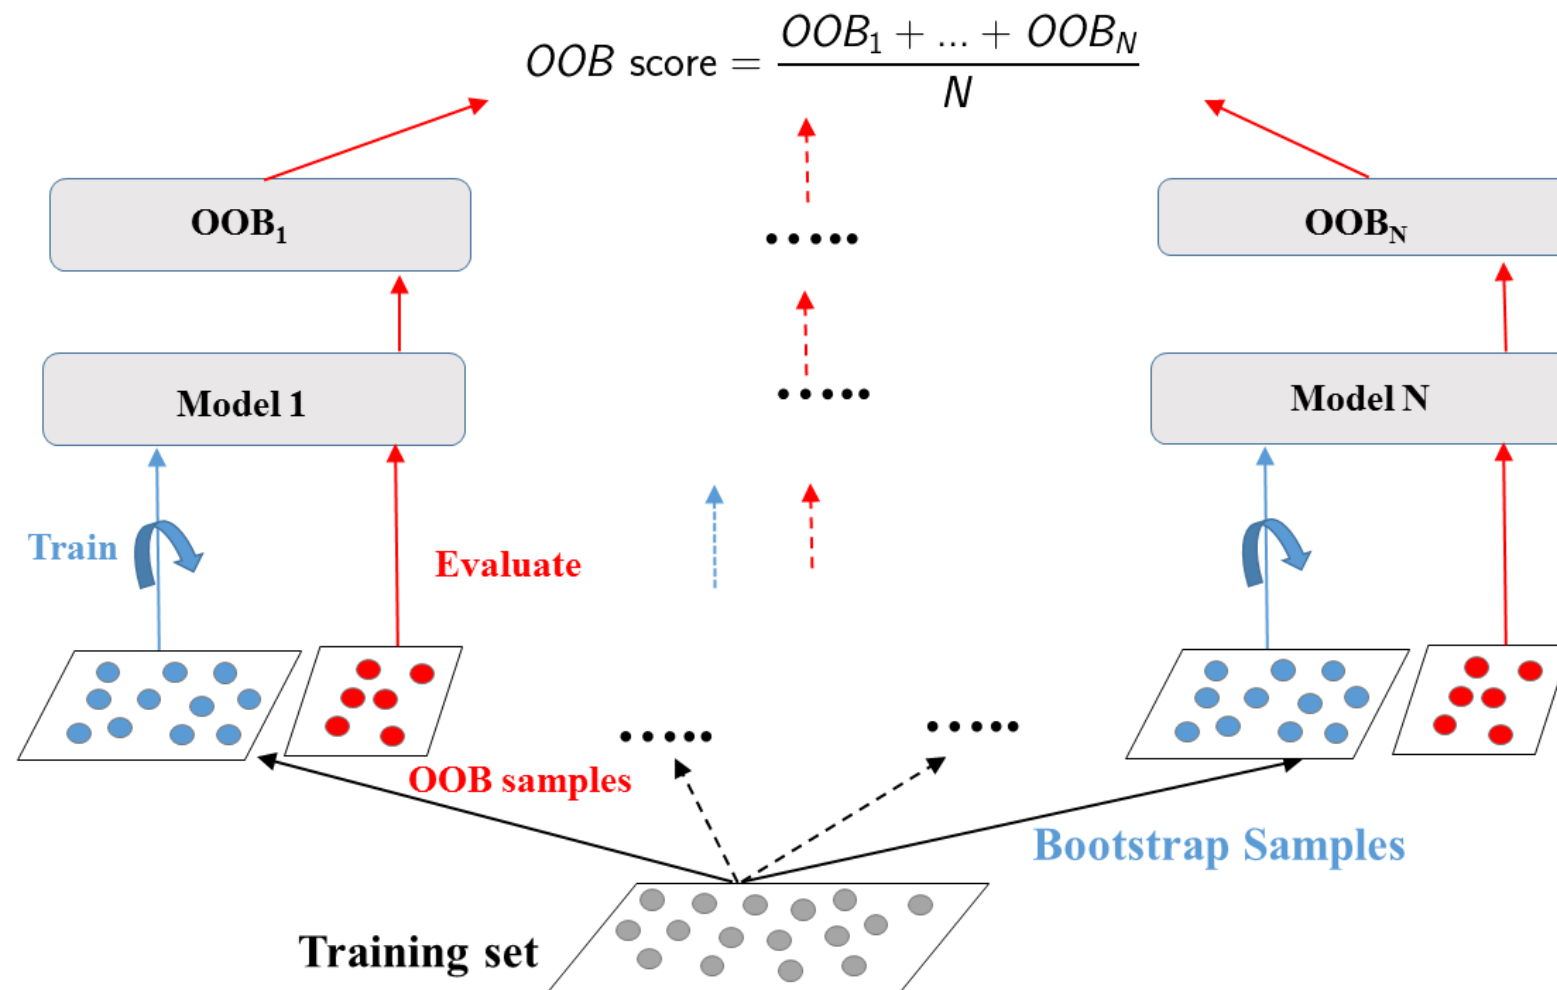


Out Of Bag (OOB) instances

- On average, for each model, 63% of the training instances are sampled.
- The remaining 37% constitute the OOB instances.



OOB Evaluation



OOB Evaluation in sklearn (Breast Cancer Dataset)

```
# Import models and split utility function
In [1]: from sklearn.ensemble import BaggingClassifier
In [2]: from sklearn.tree import DecisionTreeClassifier
In [3]: from sklearn.metrics import accuracy_score
In [4]: from sklearn.model_selection import train_test_split

# Set seed for reproducibility
In [5]: SEED = 1

# Split data into 70% train and 30% test
In [6]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size= 0.3,
                        stratify= y,
                        random_state=SEED)
```



OOB Evaluation in sklearn (Breast Cancer Dataset)

```
# Instantiate a classification-tree 'dt'
In [7]: dt = DecisionTreeClassifier(max_depth=4,
                                   min_samples_leaf=0.16,
                                   random_state=SEED)

# Instantiate a BaggingClassifier 'bc'; set oob_score= True
In [8]: bc = BaggingClassifier(base_estimator=dt, n_estimators=300,
                               oob_score=True, n_jobs=-1)

# Fit 'bc' to the training set
In [9]: bc.fit(X_train, y_train)

# Predict the test set labels
In [10]: y_pred = bc.predict(X_test)
```




OOB Evaluation in sklearn (Breast Cancer Dataset)

```
# Evaluate test set accuracy
In [11]: test_accuracy = accuracy_score(y_test, y_pred)

# Extract the OOB accuracy from 'bc'
In [12]: oob_accuracy = bc.oob_score_

# Print test set accuracy
In [13]: print('Test set accuracy: {:.3f}'.format(test_accuracy))
Out[13]: Test set accuracy: 0.936

# Print OOB accuracy
In [14]: print('OOB accuracy: {:.3f}'.format(oob_accuracy))
Out[14]: OOB accuracy: 0.925
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Random Forests

Elie Kawerk
Data Scientist



Bagging

- Base estimator: Decision Tree, Logistic Regression, Neural Net, ...
- Each estimator is trained on a distinct bootstrap sample of the training set
- Estimators use all features for training and prediction

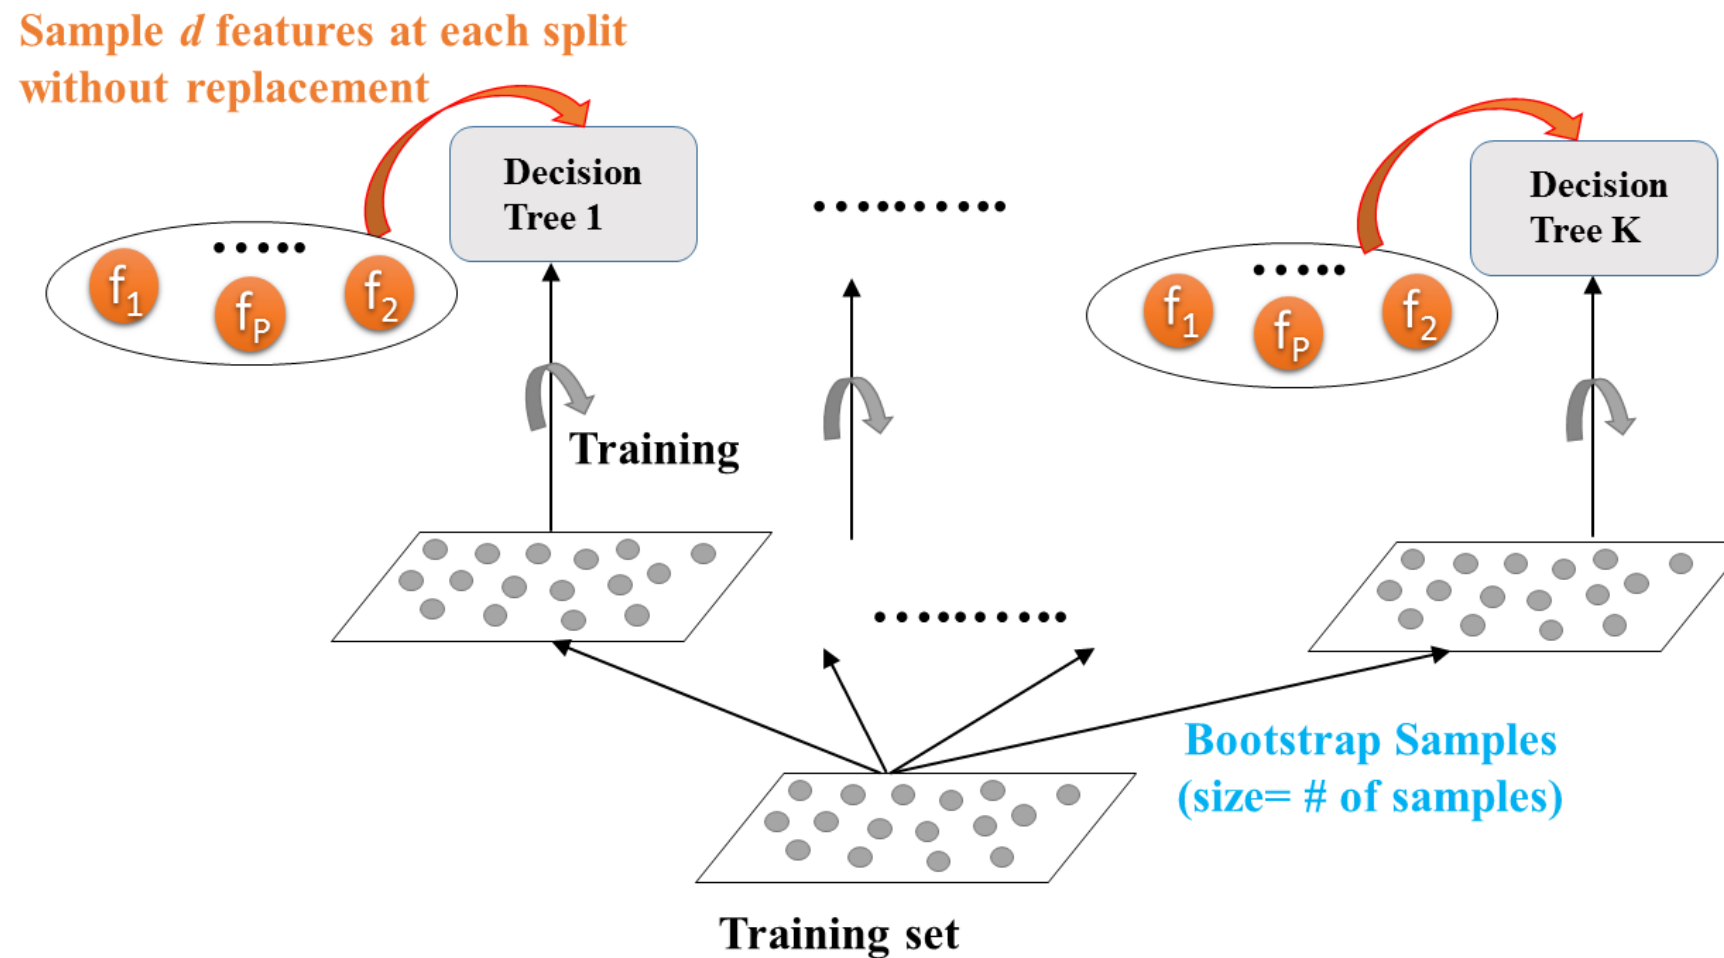


Further Diversity with Random Forests

- Base estimator: Decision Tree
- Each estimator is trained on a different bootstrap sample having the same size as the training set
- RF introduces further randomization in the training of individual trees
- d features are sampled at each node without replacement
($d < \text{total number of features}$)

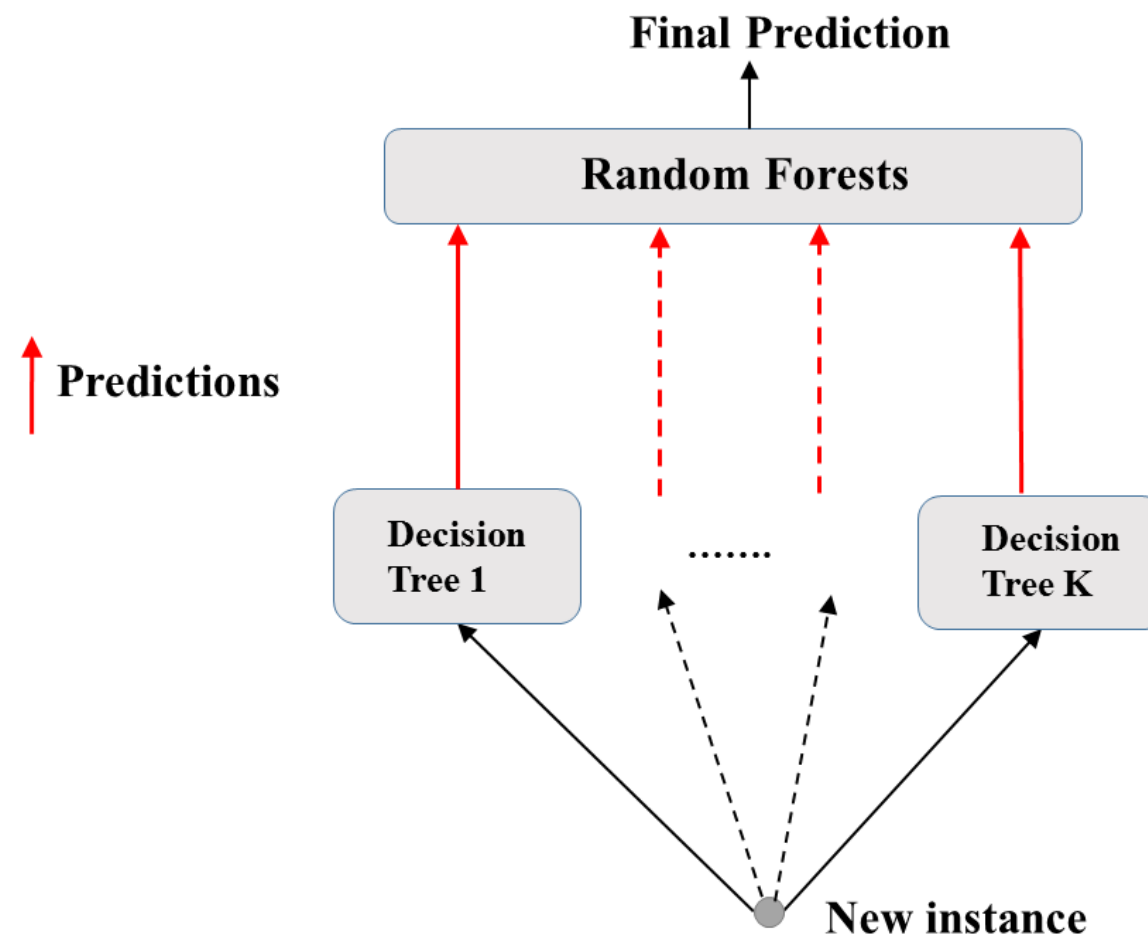


Random Forests: Training





Random Forests: Prediction





Random Forests: Classification & Regression

Classification:

- Aggregates predictions by majority voting
- `RandomForestClassifier` in `scikit-learn`

Regression:

- Aggregates predictions through averaging
- `RandomForestRegressor` in `scikit-learn`

Random Forests Regressor in sklearn (auto dataset)

```
# Basic imports
In [1]: from sklearn.ensemble import RandomForestRegressor
In [2]: from sklearn.model_selection import train_test_split
In [3]: from sklearn.metrics import mean_squared_error as MSE

# Set seed for reproducibility
In [4]: SEED = 1

# Split dataset into 70% train and 30% test
In [5]: X_train, X_test, y_train, y_test = \
        train_test_split(X, y,
                        test_size=0.3,
                        random_state=SEED)
```



Random Forests Regressor in sklearn (auto dataset)

```
# Instantiate a random forests regressor 'rf' 400 estimators
In [6]: rf = RandomForestRegressor(n_estimators=400,
                                   min_samples_leaf=0.12,
                                   random_state=SEED)

# Fit 'rf' to the training set
In [7]: rf.fit(X_train, y_train)

# Predict the test set labels 'y_pred'
In [8]: y_pred = rf.predict(X_test)

# Evaluate the test set RMSE
In [9]: rmse_test = MSE(y_test, y_pred)**(1/2)

# Print the test set RMSE
In [10]: print('Test set RMSE of rf: {:.2f}'.format(rmse_test))

Out[10]: Test set RMSE of rf: 3.98
```



Feature Importance

Tree-based methods: enable measuring the importance of each feature in prediction.

In `sklearn`:

- how much the tree nodes use a particular feature (weighted average) to reduce impurity
- accessed using the attribute `feature_importance_`

Feature Importance in sklearn

```
In [11]: import pandas as pd
```

```
In [12]: import matplotlib.pyplot as plt
```

```
# Create a pd.Series of features importances
```

```
In [13]: importances_rf = pd.Series(rf.feature_importances_,
                                     index = X.columns)
```

```
# Sort importances rf
```


```
In [14]: sorted importances rf = importances rf.sort values()
```

```
# Make a horizontal bar plot
```

[illegible]



Feature Importance in sklearn





MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!