



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Tuning a CART's hyperparameters

Elie Kawerk

Data Scientist



Hyperparameters

Machine learning model:

- **parameters**: learned from data
 - CART example: split-point of a node, split-feature of a node, ...
- **hyperparameters**: not learned from data, set prior to training
 - CART example: `max_depth`, `min_samples_leaf`, **splitting criterion** ...



What is hyperparameter tuning?

- **Problem:** search for a set of optimal hyperparameters for a learning algorithm.
- **Solution:** find a set of optimal hyperparameters that results in an optimal model.
- **Optimal model:** yields an optimal **score**.
- **Score:** in sklearn defaults to accuracy (classification) and R^2 (regression).
- Cross validation is used to estimate the generalization performance.



Why tune hyperparameters?

- In `sklearn`, a model's default hyperparameters are not optimal for all problems.
- Hyperparameters should be tuned to obtain the best model performance.



Approaches to hyperparameter tuning

- Grid Search
- Random Search
- Bayesian Optimization
- Genetic Algorithms
-



Grid search cross validation

- Manually set a grid of discrete hyperparameter values.
- Set a metric for scoring model performance.
- Search exhaustively through the grid.
- For each set of hyperparameters, evaluate each model's CV score.
- The optimal hyperparameters are those of the model achieving the best CV score.

Grid search cross validation: example

- Hyperparameters grids:
 - `max_depth = {2,3,4},`
 - `min_samples_leaf = {0.05, 0.1}`
- hyperparameter space = $\{ (2,0.05) , (2,0.1) , (3,0.05), \dots \}$
- CV scores = $\{ score_{(2,0.05)} , \dots \}$
- optimal hyperparameters = set of hyperparameters corresponding to the best CV score.



Inspecting the hyperparameters of a CART in sklearn

```
# Import DecisionTreeClassifier
In [1]: from sklearn.tree import DecisionTreeClassifier

# Set seed to 1 for reproducibility
In [2]: SEED = 1

# Instantiate a DecisionTreeClassifier 'dt'
In [3]: dt = DecisionTreeClassifier(random_state=SEED)
```




Inspecting the hyperparameters of a CART in sklearn

```
# Print out 'dt's hyperparameters  
In [4]: print(dt.get_params())
```

Out[4]:

```
{ 'class_weight': None,  
  'criterion': 'gini',  
  'max_depth': None,  
  'max_features': None,  
  'max_leaf_nodes': None,  
  'min_impurity_decrease': 0.0,  
  'min_impurity_split': None,  
  'min_samples_leaf': 1,  
  'min_samples_split': 2,  
  'min_weight_fraction_leaf': 0.0,  
  'presort': False,  
  'random_state': 1,  
  'splitter': 'best' }
```

Grid search CV in sklearn (Breast Cancer dataset)

```
# Import GridSearchCV
In [5]: from sklearn.model_selection import GridSearchCV

# Define the grid of hyperparameters 'params_dt'
In [6]: params_dt = {
        'max_depth': [3, 4, 5, 6],
        'min_samples_leaf': [0.04, 0.06, 0.08],
        'max_features': [0.2, 0.4, 0.6, 0.8]
        }

# Instantiate a 10-fold CV grid search object 'grid_dt'
In [7]: grid_dt = GridSearchCV(estimator=dt,
                               param_grid=params_dt,
                               scoring='accuracy',
                               cv=10,
                               n_jobs=-1)

# Fit 'grid_dt' to the training data
In [8]: grid_dt.fit(X_train, y_train)
```



Extracting the best hyperparameters

```
# Extract best hyperparameters from 'grid_dt'
In [9]: best_hyperparams = grid_dt.best_params_

In [10]: print('Best hyperparameters:\n', best_hyperparams)

Out[10]: Best hyperparameters:
         {'max_depth': 3, 'max_features': 0.4, 'min_samples_leaf': 0.06}

# Extract best CV score from 'grid_dt'
In [11]: best_cv_score = grid_dt.best_score_

In [12]: print('Best CV accuracy'.format(best_cv_score))

Out[12]: Best CV accuracy: 0.938
```



Extracting the best estimator

```
# Extract best model from 'grid_dt'
In [13]: best_model = grid_dt.best_estimator_

# Evaluate test set accuracy
In [14]: test_acc = best_model.score(X_test, y_test)

# Print test set accuracy
In [15]: print("Test set accuracy of best model: {:.3f}".format(test_acc))

Out[15]: Test set accuracy of best model: 0.947
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Tuning an RF's Hyperparameters

Elie Kawerk

Data Scientist



Random Forests Hyperparameters

- CART hyperparameters
- number of estimators
- bootstrap
-



Tuning is expensive

Hyperparameter tuning:

- computationally expensive,
- sometimes leads to very slight improvement,

Weight the impact of tuning on the whole project.

Inspecting RF Hyperparameters in sklearn

```
# Import RandomForestRegressor
In [1]: from sklearn.ensemble import RandomForestRegressor

# Set seed for reproducibility
In [2]: SEED = 1

# Instantiate a random forests regressor 'rf'
In [3]: rf = RandomForestRegressor(random_state= SEED)
```



Inspecting RF Hyperparameters in sklearn

```
# Inspect rf's hyperparameters
In [4]: rf.get_params()

Out[4]: {'bootstrap': True,
        'criterion': 'mse',
        'max_depth': None,
        'max_features': 'auto',
        'max_leaf_nodes': None,
        'min_impurity_decrease': 0.0,
        'min_impurity_split': None,
        'min_samples_leaf': 1,
        'min_samples_split': 2,
        'min_weight_fraction_leaf': 0.0,
        'n_estimators': 10,
        'n_jobs': -1,
        'oob_score': False,
        'random_state': 1,
        'verbose': 0,
        'warm_start': False}
```

GridSearchCV in sklearn (auto dataset)

```
# Basic imports
In [5]: from sklearn.metrics import mean_squared_error as MSE
In [6]: from sklearn.model_selection import GridSearchCV

# Define a grid of hyperparameter 'params_rf'
In [7]: params_rf = {
        'n_estimators': [300, 400, 500],
        'max_depth': [4, 6, 8],
        'min_samples_leaf': [0.1, 0.2],
        'max_features': ['log2', 'sqrt']
    }

# Instantiate 'grid_rf'
In [8]: grid_rf = GridSearchCV(estimator=rf,
                               param_grid=params_rf,
                               cv=3,
                               scoring='neg_mean_squared_error',
                               verbose=1,
                               n_jobs=-1)
```

Searching for the best hyperparameters

```
# Fit 'grid_rf' to the training set
In [9]: grid_rf.fit(X_train, y_train)
```

Out[9]:

```
Fitting 3 folds for each of 36 candidates, totalling 108 fits
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed: 10.0s
[Parallel(n_jobs=-1)]: Done 108 out of 108 | elapsed: 24.3s finished
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=4,
                      max_features='log2', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=0.1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=400, n_jobs=1,
                      oob_score=False, random_state=1, verbose=0, warm_start=False)
```



Extracting the best hyperparameters

```
# Extract best hyperparameters from 'grid_rf'
In [10]: best_hyperparams = grid_rf.best_params_

In [11]: print('Best hyperparameters:\n', best_hyperparams)

Out[11]: Best hyperparameters:
        {'max_depth': 4,
         'max_features': 'log2',
         'min_samples_leaf': 0.1,
         'n_estimators': 400}
```



Evaluating the best model performance

```
# Extract best model from 'grid_rf'
In [12]: best_model = grid_rf.best_estimator_

# Predict the test set labels
In [13]: y_pred = best_model.predict(X_test)

# Evaluate the test set RMSE
In [14]: rmse_test = MSE(y_test, y_pred)**(1/2)

# Print the test set RMSE
In [15]: print('Test set RMSE of rf: {:.2f}'.format(rmse_test))

Out[15]: Test set RMSE of rf: 3.89
```



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Let's practice!



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Congratulations!

Elie Kawerk
Data Scientist



How far you have come

- **Chapter 1:** Decision-Tree Learning
- **Chapter 2:** Generalization Error, Cross-Validation, Ensembling
- **Chapter 3:** Bagging and Random Forests
- **Chapter 4:** AdaBoost and Gradient-Boosting
- **Chapter 5:** Model Tuning



MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Thank you!